# Oracle Linux 8 Monitoring and Tuning the System





Oracle Linux 8 Monitoring and Tuning the System,

F24025-12

Copyright  $\ensuremath{\texttt{@}}$  2019, 2022, Oracle and/or its affiliates.

### Contents

#### Preface

Conventions	\
Documentation Accessibility	\
Access to Oracle Support for Accessibility	\
Diversity and Inclusion	\
Monitoring the System and Optimizing Performance	ce
Working With System Performance and Monitoring Utilities	1-
Monitoring the Usage of System Resources	1-7
Monitoring CPU Usage	1-3
Monitoring Memory Usage	1
Using Adaptive Memory Management	1
Monitoring Block I/O Usage	1
Monitoring File System Usage	1-
Monitoring Network Usage	1-
Using the Graphical System Monitor	1-
Working With the sos Command	
Installing the sos Package	2-:
Running the sos Command	2-:
Reviewing Information Gathered by sosreport	2-
Working With OSWatcher Black Box	
Installing OSWbb	3-:
Running OSWbb	3-
Analyzing OSWbb Archived Files	3-:
Working With Performance Co-Pilot	
Installing PCP	4-:



Reviewing Information Gathered by PCP	4-1
Using PCP Monitor Host to Analyze Performance Metrics	4-2
Review Live Performance Metrics in Real Time	4-2
Review Recorded Performance Metrics	4-2
Review Details About Recorded Performance Metrics	4-2
Validate System Status When Performance Metrics Were Captured	4-3
Working With Tuned	
About Tuned Profiles	5-1
About the Default Tuned Profiles	5-2
About Static and Dynamic Tuning in Tuned	5-3
Installing and Enabling Tuned by Using the Command Line	5-3
Running Tuned in no-daemon Mode	5-4
Administering the Tuned Service and Tuned Profiles	5-4
Listing Tuned Profiles	5-4
Activating a Tuned Profile	5-5
Disabling Tuned	5-5
Automating System Tasks	
Configuring cron Jobs	6-1
Controlling Access to Running cron Jobs	6-3
Configuring anacron Jobs	6-3
Running One-Time Tasks	6-4
Changing the Behavior of Batch Jobs	6-5
Configuring and Using Auditing	
Working With System Log files	7-2
About Logging Configuration (/etc/rsyslog.conf)	7-2
Configuring Logwatch	7-6
Using Process Accounting	7-6
Working With Kernel Dumps	
Kdump Installation and Configuration	8-1
Files That Are Used by Kdump	8-1
Installing and Configuring Kdump	8-2
Installing and Configuring Kdump Configuring the Kdump Output Location	8-2 8-3



Using Early Kdump	8-4
Using Kdump with OCFS2	8-5



#### **Preface**

Oracle® Linux 8: Monitoring and Tuning the System describes the various utilities, features, and services that you can use to monitor system performance, detect performance issues, and improve the performance of various system components.

#### Conventions

The following text conventions are used in this document:

Convention	Meaning
boldface	Boldface type indicates graphical user interface elements associated with an action, or terms defined in text or the glossary.
italic	Italic type indicates book titles, emphasis, or placeholder variables for which you supply particular values.
monospace	Monospace type indicates commands within a paragraph, URLs, code in examples, text that appears on the screen, or text that you enter.

### **Documentation Accessibility**

For information about Oracle's commitment to accessibility, visit the Oracle Accessibility Program website at https://www.oracle.com/corporate/accessibility/.

For information about the accessibility of the Oracle Help Center, see the Oracle Accessibility Conformance Report at https://www.oracle.com/corporate/accessibility/templates/t2-11535.html.

### Access to Oracle Support for Accessibility

Oracle customers that have purchased support have access to electronic support through My Oracle Support. For information, visit <a href="https://www.oracle.com/corporate/accessibility/learning-support.html#support-tab">https://www.oracle.com/corporate/accessibility/learning-support.html#support-tab</a>.

### **Diversity and Inclusion**

Oracle is fully committed to diversity and inclusion. Oracle respects and values having a diverse workforce that increases thought leadership and innovation. As part of our initiative to build a more inclusive culture that positively impacts our employees, customers, and partners, we are working to remove insensitive terms from our



products and documentation. We are also mindful of the necessity to maintain compatibility with our customers' existing technologies and the need to ensure continuity of service as Oracle's offerings and industry standards evolve. Because of these technical constraints, our effort to remove insensitive terms is ongoing and will take time and external cooperation.



1

## Monitoring the System and Optimizing Performance

Performance issues can be caused by a number of your system's components, including software or hardware, as well as any related interactions. Many performance diagnostics utilities are available in Oracle Linux and include tools that monitor and analyze the resource usage of different hardware components, as well as tracing tools for diagnosing performance issues in multiple processes and related threads.

Many performance issues are the result of configuration errors. You can avoid these errors by using a validated configuration that has been pre-tested for the supported software, hardware, storage, drivers, and networking components. A validated configuration incorporates best practices for an Oracle Linux deployment and has undergone real-world testing of the complete stack. Oracle publishes many validated configurations, which are freely available for download. Refer to the release notes for the release that you are running for additional recommendations on kernel parameter settings.

### Working With System Performance and Monitoring Utilities

The following utilities enable you to collect information about system resource usage and errors, and help you to identify performance problems that are caused by overloaded disks, network, memory, or CPUs:

#### dmesg

Displays the contents of the kernel ring buffer, which can contain errors about system resource usage. Provided by the util-linux package.

#### dstat

Displays statistics about system resource usage. Provided by the dstat package.

#### free

Displays the amount of free and used memory in the system. Provided by the procps package, which is install by default in Oracle Linux 8.

#### iostat

Reports I/O statistics. Provided by the sysstat package.

#### iotop

Monitors disk and swap I/O on a per-process basis. Provided by the iotop package.

#### ip

Reports network interface statistics and errors. Provided by the iproute package, which is installed by default in Oracle Linux 8.

#### mpstat

Reports processor-related statistics. Provided by the sysstat package.



#### nfsiostat

Reports I/O statistics for NFS mounts. Provided by the nfs-utils package.

#### sar

Reports information about system activity. Provided by the sysstat package.

#### SS

Reports network interface statistics. Provided by the iproute package.

#### top

Provides a dynamic real-time view of the tasks that are running on a system. Provided by the procps package.

#### uptime

Displays the system load averages for the past 1, 5, and 15 minutes. Provided by the procps package.

#### vmstat

Reports virtual memory statistics. Provided by the procps package, which is installed by default in Oracle Linux 8.

Many of these utilities provide overlapping functionality. For more information, see the individual manual page for the utility.

See https://docs.oracle.com/en/learn/system\_monitoring\_linux8/ for a hands-on tutorial and associated video content on many of these utilities.

### Monitoring the Usage of System Resources

You need to collect and monitor system resources so that you are provided with a continuous record of a system's performance. First, establish a baseline of acceptable measurements under typical operating conditions. You can then use that baseline as a reference point to make it easier to identify memory shortages, spikes in resource usage, and other problems when they occur. Monitoring system performance also enables you to plan for future growth and determine how configuration changes might affect future performance.

To run a monitoring command for a set number of seconds in real time and watch the output change, use the watch command. For example, run the mpstat command once per second with the following command:

```
sudo watch -n 1 mpstat
```

Alternatively, many of the commands enable you to specify the sampling interval in seconds, for example:

```
sudo mpstat seconds
```

If it is installed, the sar command records statistics every 10 minutes while the system is running and retains that information for every day of the current month. The following command displays all the statistics that sar recorded for the day (DD) of the current month:

```
sudo sar -A -f /var/log/sa/saDD
```

To run the sar command as a background process and collect data in a file that you can display later by using the -f option:



sudo sar -o datafile seconds count >/dev/null 2>&1 &

In the previous command, *count* is the number of samples to record.

See also Working With OSWatcher Black Box and Working With Performance Co-Pilot.

### Monitoring CPU Usage

The uptime, mpstat, sar, dstat, and top utilities enable you to monitor CPU usage. When your system's CPU cores are occupied with executing code processes, other processes must wait until a CPU core becomes free or when the scheduler switches a CPU core to run its code. If too many processes are queued too often, then that can represent a bottleneck in the performance of the system.

The commands mpstat -P ALL and sar -u -P ALL display CPU usage statistics for each CPU core and is averaged across all of the CPU cores.

The <code>%idle</code> value shows the percentage of time that a CPU was not running system or process code. If the value of <code>%idle</code> is near 0% most of the time on all CPU cores, the system is CPU-bound for the workload that it is running. The percentage of time spent running system code (<code>%system</code> or <code>%sys</code>) should not usually exceed 30%, especially if <code>%idle</code> is close to 0%.

The system load average represents the number of processes that are running on CPU cores, waiting to run, or waiting for disk I/O activity to complete averaged over a period of time. On a busy system, the load average reported by  ${\tt uptime}$  or  ${\tt sar}$   $-{\tt q}$  should usually be not greater than two times the number of CPU cores over periods as long as 5 or 15 minutes. If the load average exceeds four times the number of CPU cores for long periods, the system is overloaded.

In addition to load averages (ldavg-\*), the sar-q command reports the number of processes currently waiting to run (the *run-queue size*, runq-sz) and the total number of processes ( $plist\ sz$ ). The value of runq-sz also provides an indication of CPU saturation.

Determine the system's average load under normal loads, where users and applications do not experience problems with system responsiveness, and then look for deviations from this benchmark over time. A dramatic rise in the load average can indicate a serious performance problem.

A combination of sustained large load average or large run queue size and low <code>%idle</code> can indicate that the system has insufficient CPU capacity for the workload. When CPU usage is high, use a command such as <code>dstat</code> or <code>top</code> to determine which processes are most likely to be responsible. For example, the following <code>dstat</code> command shows which processes are using CPUs, memory, and block I/O most intensively:

```
dstat --top-cpu --top-mem --top-bio
```

The top command provides a real-time display of CPU activity. By default, top lists the most CPU-intensive processes on the system. In its upper section, top displays general information including the load averages over the past 1, 5 and 15 minutes, the number of running and sleeping processes (tasks), and total CPU and memory usage. In its lower section, top displays a list of processes, including the process ID number (PID), the process owner, CPU usage, memory usage, running time, and the command name. By default, the list is sorted by CPU usage, with the top consumer of CPU listed first. Type f to select which fields top displays, o to change the order of the fields, or o to change the sort field. For example, entering on sorts the list on the percentage memory usage field (%MEM).



### Monitoring Memory Usage

The sar -r command reports memory utilization statistics, including %memused, which is the percentage of physical memory in use.

sar —B reports memory paging statistics, including pgscank/s, which is the number of memory pages scanned by the kswapd daemon per second, and pgscand/s, which is the number of memory pages scanned directly per second.

sar -W reports swapping statistics, including pswpin/s and pswpout/s, which are the numbers of pages per second swapped in and out per second.

If %memused is near 100% and the scan rate is continuously over 200 pages per second, the system has a memory shortage.

Once a system runs out of real or physical memory and starts using swap space, its performance deteriorates dramatically. If you run out of swap space, your programs or the entire operating system are likely to crash. If free or top indicate that little swap space remains available, this is also an indication you are running low on memory.

The output from the <code>dmesg</code> command might include notification of any problems with physical memory that were detected at boot time.

#### Using Adaptive Memory Management

The Adaptive Memory Management daemon is a user space service that monitors free memory on your Oracle Linux system and predicts memory fragmentation and usage. It can also automatically reclaim memory if the system if memory becomes too fragmented or is at risk of being filled to capacity.

If the system memory becomes highly fragmented, adaptivemmd triggers the kernel to compact memory so that fragmented space can be reclaimed before it is reallocated. If the system is likely to exhaust the available memory, watermarks are adjusted and this can trigger the kernel to free up new pages in memory. Adaptive Memory Management is available in Unbreakable Enterprise Kernel Release 6 and later.

Use dnf to install the adaptivemm package from the ol8\_UEKR6 yum repository or equivalent ULN channel:

```
sudo dnf install -y adaptivemm
```

To enable and run the adaptivemmd service:

```
sudo systemctl enable --now adaptivemmd
```

You can change the configuration options in /etc/sysconfig/adaptivemmd.

For more information see the adaptivemmd (8) manual page.

### Monitoring Block I/O Usage

The iostat command monitors the loading of block I/O devices by observing the time that the devices are active relative to the average data transfer rates. You can use this information to adjust the system configuration to balance the I/O loading across disks and host adapters.



iostat -x reports extended statistics about block I/O activity at one second intervals, including util, which is the percentage of CPU time spent handling I/O requests to a device, and avgqu-sz, which is the average queue length of I/O requests that were issued to that device. If util approaches 100% or avgqu-sz is greater than 1, device saturation is occurring.

You can also use the sar -d command to report on block I/O activity, including values for to subseteq the sar -d command to report on block I/O activity, including values for to subseteq to subseteq the sar -d command to report on block I/O activity, including values for to subseteq to

The iotop utility can help you identify which processes are responsible for excessive disk I/O. iotop has a similar user interface to top. In its upper section, iotop displays the total disk input and output usage in bytes per second. In its lower section, iotop displays I/O information for each process, including disk input output usage in bytes per second, the percentage of time spent swapping in pages from disk or waiting on I/O, and the command name. Use the left and right arrow keys to change the sort field, and press A to toggle the I/O units between bytes per second and total number of bytes, or  $\circ$  to toggle between displaying all processes or only those processes that are performing I/O.

### Monitoring File System Usage

The sar -v command reports the number of unused cache entries in the directory cache (dentunusd) and the numbers of in-use file handles (file-nr), inode handlers (inode-nr), and pseudo terminals (pty-nr).

nfsiostat reports I/O statistics for each NFS file system that is mounted. If this command is not available install the nfs-utils package.

### Monitoring Network Usage

The ip -s link command displays network statistics and errors for all network devices, including the numbers of bytes transmitted (TX) and received (RX). The dropped and overrun fields provide an indicator of network interface saturation, for example:

```
ip -s link
```

#### The following output is displayed:

```
1: lo: <LOOPBACK, UP, LOWER UP> mtu 65536 qdisc noqueue state UNKNOWN mode DEFAULT group
default qlen 1000
   link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00
   RX: bytes packets errors dropped overrun mcast
   240 4 0 0 0
   TX: bytes packets errors dropped carrier collsns
   240 4 0
                          0
                                0
                                         0
2: enp0s3: <BROADCAST,MULTICAST,UP,LOWER UP> mtu 1500 qdisc fq codel state UP mode
DEFAULT group default glen 1000
   link/ether 08:00:27:60:95:d5 brd ff:ff:ff:ff:ff
   RX: bytes packets errors dropped overrun mcast
   258187485 671730 0
                         0
                                 0
   TX: bytes packets errors dropped carrier collsns
   13227598 130827 0
                           0
                                   0
```

The ss -s command displays summary statistics for each protocol, for example:

ss -s



#### The following output is displayed:

```
Total: 193
TCP: 9 (estab 2, closed 0, orphaned 0, timewait 0)
```

### Using the Graphical System Monitor

The GNOME desktop environment includes a graphical system monitor that enables you to display information about the system configuration, running processes, resource usage, and file systems.

To display the System Monitor, use the following command:

```
gnome-system-monitor
```

Selecting the **Resources** tab displays the following information:

- CPU usage history in graphical form and the current CPU usage as a percentage.
- Memory and swap usage history in graphical form and the current memory and swap usage.
- Network usage history in graphical form, the current network usage for reception and transmission, and the total amount of data received and transmitted.

To display the System Monitor Manual, press F1 or select Help, then select Contents.



### Working With the sos Command

The sos command collects information about a system such as hardware configuration, software configuration, and operational state. You can also use the sos report command to enable diagnostics and analytical functions on the current system. To assist in troubleshooting a problem, the sos report command records the information in a compressed file that you can send to a support representative.

### Installing the sos Package

If the sos package is not already installed on your system, use dnf to install it:

```
sudo dnf install sos
```

Use the sos report -1 command to list the available plugins and plugin options:

See the sos-report (1) manual page for information about how to enable or disable plugins and how to set values for plugin options.

### Running the sos Command

You can run the sos report command to record information about a problem area and specify options to tailor the report it generates as follows:

```
sudo sos report [options ...]
```

For example, to record only information about Apache and Tomcat and to gather all of the Apache logs, use the following command:

```
sudo sos report -o apache, tomcat -k apache.log=on
```

The output is similar to the following:

```
sosreport (version 4.1)
.
.
.
Press ENTER to continue, or CTRL-C to quit.
```

To enable all of the boolean options for all of the loaded plugins (excluding the  ${\tt rpm.rpmva}$  plugin) and verify all packages:

```
sudo sos report -a -k rpm.rpmva=off
```

Note that this process can take a considerable amount of time, but once it has completed, press Enter and then provide any additional information that is required:

The  ${\tt sos}$  report command saves the report as an  ${\tt xz}$ -compressed  ${\tt tar}$  file in the  ${\tt /tmp}$  directory.

Optionally, to obfuscate sensitive information, you can run the sos clean command on the compressed archive generated from the sos report command.

The cleaned report obfuscates the following details:

- IPv4 addresses and networks (network topologies are retained)
- MAC addresses
- Host names
- User names
- Any words or phrases that you specify with the --keyword option





Reports processed with the sos clean command obfuscate certain details that may be needed for advanced troubleshooting, such as networking information.

To generate a cleaned report, run the sos clean command on the compressed archive generated from the sos report command in the /tmp directory:

```
sudo sos clean /var/tmp/sosreport-hostname-case#-datestamp-ID.tar.xz
```

Press Enter to proceed. After the sos clean command completes, a new xz-compressed tar file with -obfuscated in the file name is created the /tmp directory.

You can also perform this step automatically by appending the --clean option to the sos report command, as follows:

```
sudo sos report --clean [options ...]
```

For more information, see the <code>sos-report(1)</code> and <code>sos-clean(1)</code> manual pages. See also https://github.com/sosreport/sos/wiki.

### Reviewing Information Gathered by sosreport

The  ${\tt sos}$  command is automatically configured to collect hardware information, system configuration files and log data; but, you can enable and disable modules to suit your own data protection needs.



The module information that is provided in this table relates to sos 3.9. To verify the modules you have installed on your system, read Installing the sos Package.

Disabling modules prevents the sos command from collecting certain details that may be needed for advanced troubleshooting, such as networking information.

Module	Information Type	Included Files
anaconda	Installation log files	<ul><li>/root/install.log</li><li>/root/ install.log.syslog</li><li>/var/log/anaconda</li></ul>
auditd	Audit log files	<ul><li>/var/log/anaconda.*</li><li>/etc/audit/auditd.conf</li><li>/etc/audit/audit.rules</li><li>/var/log/audit/*</li></ul>



Module	Information Type	Included Files
boot	System boot process details	<ul> <li>/etc/milo.conf</li> <li>/etc/silo.conf</li> <li>/boot/efi/efi/redhat/elilo.conf</li> <li>/etc/yaboot.conf</li> <li>/boot/yaboot.conf</li> </ul>
cron	Root user cron commands	<ul><li>/etc/cron*</li><li>/etc/crontab</li><li>/var/log/cron</li><li>/var/spool/cron</li></ul>
cups	Printer log files	<ul><li>/etc/cups/*.conf</li><li>/etc/cups/*.types</li><li>/etc/cups/lpoptions</li><li>/etc/cups/ppd/*.ppd</li><li>/var/log/cups/*</li></ul>
date	Context data	<ul> <li>/etc/localtime</li> </ul>
devicemapper	Hardware details	
filesys	List of all files in use	<ul> <li>/proc/fs/*</li> <li>/proc/mounts</li> <li>/proc/filesystems</li> <li>/proc/self/mounts</li> <li>/proc/self/mountinfo</li> <li>/proc/self/mountstats</li> <li>/proc/[0-9]*/mountinfo</li> <li>/etc/mtab</li> <li>/etc/fstab</li> </ul>
grub2	Kernel and system start-up configuration	<ul> <li>/boot/efi/EFI/*/ grub.cfg</li> <li>/boot/grub2/grub.cfg</li> <li>/boot/grub/grub.cfg</li> <li>/boot/grub/grub.cfg</li> <li>/boot/loader/entries</li> <li>/etc/default/grub</li> <li>/etc/grub2.cfg</li> <li>/etc/grub.d/*</li> </ul>
hardware	Hardware details	<ul> <li>/proc/interrupts</li> <li>/proc/irq</li> <li>/proc/dma</li> <li>/proc/devices</li> <li>/proc/rtc</li> <li>/var/log/mcelog</li> <li>/sys/class/dmi/id/*</li> <li>/sys/class/drm/*/edid</li> </ul>



Module	Information Type	Included Files
host	Host identification	<ul><li>/etc/sos.conf</li></ul>
		<ul><li>/etc/hostid</li></ul>



Module	Information Type	Included Files
kernel	System log files	• /etc/conf.modules
		<ul><li>/etc/modules.conf</li></ul>
		<ul><li>/etc/modprobe.conf</li></ul>
		<ul><li>/etc/modprobe.d</li></ul>
		<ul><li>/etc/sysctl.conf</li></ul>
		<ul><li>/etc/sysctl.d</li></ul>
		<ul><li>/lib/modules/*/</li></ul>
		modules.dep
		<ul> <li>/lib/sysctl.d</li> </ul>
		<ul><li>/proc/cmdline</li></ul>
		<ul><li>/proc/driver</li></ul>
		• /proc/kallsyms
		<pre> /proc/lock*</pre>
		<ul><li>/proc/buddyinfo</li></ul>
		• /proc/misc
		• /proc/modules
		• /proc/slabinfo
		• /proc/softirqs
		• /proc/sys/kernel/
		random/boot id
		<ul><li>/proc/sys/kernel/</li></ul>
		tainted
		<ul><li>/proc/timer*</li></ul>
		<ul><li>/proc/zoneinfo</li></ul>
		<ul><li>/sys/firmware/acpi/*</li></ul>
		<ul><li>/sys/kernel/debug/</li></ul>
		tracing/*
		<ul><li>/sys/kernel/</li></ul>
		livepatch/*
		<ul><li>/sys/module/*/</li></ul>
		parameters
		<ul><li>/sys/module/*/</li></ul>
		initstate
		<ul><li>/sys/module/*/refcnt</li></ul>
		<ul><li>/sys/module/*/taint</li></ul>
		<ul><li>/sys/module/*/version</li></ul>
		<ul><li>/sys/devices/system/</li></ul>
		clocksource/*/
		available_clocksource
		<ul> <li>/sys/devices/system/</li> </ul>
		clocksource/*/
		current_clocksource
		• /sys/fs/pstore
		<ul><li>/var/log/dmesg</li></ul>



Module	Information Type	Included Files
libraries	List of shared libraries	<ul> <li>/etc/ld.so.conf</li> </ul>
		<ul><li>/etc/ld.so.conf.d/*</li></ul>
logs	System log files	<ul><li>/etc/syslog.conf</li></ul>
		<ul><li>/etc/rsyslog.conf</li></ul>
		<ul><li>/etc/rsyslog.d</li></ul>
		<ul><li>/run/log/journal/*</li></ul>
		<ul><li>/var/log/auth.log</li></ul>
		<ul><li>/var/log/auth.log.1</li></ul>
		<ul><li>/var/log/auth.log.2*</li></ul>
		<ul><li>/var/log/boot.log</li></ul>
		<ul> <li>/var/log/dist-upgrade</li> </ul>
		<ul> <li>/var/log/installer</li> </ul>
		<ul><li>/var/log/journal/*</li></ul>
		<ul><li>/var/log/kern.log</li></ul>
		<ul><li>/var/log/kern.log.1</li></ul>
		<ul><li>/var/log/kern.log.2*</li></ul>
		<ul><li>/var/log/messages*</li></ul>
		<ul><li>/var/log/secure*</li></ul>
		<ul><li>/var/log/syslog</li></ul>
		<ul><li>/var/log/syslog.1</li></ul>
		<ul><li>/var/log/syslog.2*</li></ul>
		<ul><li>/var/log/udev</li></ul>
		<ul><li>/var/log/unattended- upgrades</li></ul>
lvm2	Hardware details	
memory	Hardware details	• /proc/pci
_		<ul><li>/proc/meminfo</li></ul>
		<ul><li>/proc/vmstat</li></ul>
		<ul><li>/proc/swaps</li></ul>
		<ul><li>/proc/slabinfo</li></ul>
		<ul> <li>/proc/pagetypeinfo</li> </ul>
		<ul><li>/proc/vmallocinfo</li></ul>
		<ul><li>/sys/kernel/mm/ksm</li></ul>
		<ul><li>/sys/kernel/mm/</li></ul>
		transparent_hugepage/ enabled



Module	Information Type	Included Files
networking	Network identification	• /etc/dnsmasq*
		<ul><li>/etc/host*</li></ul>
		<ul> <li>/etc/inetd.conf</li> </ul>
		<ul><li>/etc/iproute2</li></ul>
		<ul><li>/etc/network*</li></ul>
		<ul><li>/etc/nftables</li></ul>
		<ul><li>/etc/nftables.conf</li></ul>
		<ul><li>/etc/nsswitch.conf</li></ul>
		<ul><li>/etc/resolv.conf</li></ul>
		<ul><li>/etc/sysconfig/ nftables.conf</li></ul>
		<ul> <li>/etc/xinetd.conf</li> </ul>
		<ul><li>/etc/xinetd.d</li></ul>
		<ul><li>/etc/yp.conf</li></ul>
		<ul><li>/proc/net/*</li></ul>
		<ul> <li>/sys/class/net/*/ device/numa_node</li> </ul>
		<ul> <li>/sys/class/net/*/flag</li> </ul>
		<ul><li>/sys/class/net/*/ statistics/*</li></ul>
pam	Login security settings	<ul><li>/etc/pam.d/*</li></ul>
		<ul> <li>/etc/security</li> </ul>
pci	Hardware details	• /proc/bus/pci
		<ul><li>/proc/iomem</li></ul>
		<ul><li>/proc/ioports</li></ul>
process	List of all running processes	<ul> <li>/proc/sched debug</li> </ul>
	and process details	<ul><li>/proc/stat</li></ul>
		<ul><li>/proc/[0-9]*/smaps</li></ul>
processor	Hardware details	<ul> <li>/proc/cpuinfo</li> </ul>
		<ul> <li>/sys/class/cpuid</li> </ul>
		<ul> <li>/sys/devices/ system/cpu</li> </ul>
rpm	Installed software packages	<pre>var/lib/rpm/*</pre>
12	motumos contrare paenages	<ul><li>/var/log/rpmpkgs</li></ul>
sar	Resource and usage data	<ul><li>/var/log/sa/*</li></ul>
selinux	Security settings	<ul><li>/etc/sestatus.conf</li></ul>
OCILIIUA	occurry settings	<ul><li>/etc/sestatus.com</li><li>/etc/selinux</li></ul>
		<ul><li>/ecc/selinux</li><li>/var/lib/selinux</li></ul>
services	All defined system services	<ul><li>/etc/inittab</li></ul>
DOT ATCES	7 in defined system services	<ul><li>/etc/inittab</li><li>/etc/rc.d/*</li></ul>
		<ul><li>/etc/rc.local</li></ul>
ach	CCU configuration	
ssh	SSH configuration	<ul><li>/etc/ssh/ssh_config</li><li>/etc/ssh/sshd_config</li></ul>



Module	Information Type	Included Files
x11	GUI logs for the X Window System	<ul> <li>/etc/X11/*</li> <li>/var/log/Xorg.*.log</li> <li>/var/log/</li> <li>Xorg.*.log.old</li> <li>/var/log/XFree86.*.log</li> <li>/var/log/</li> </ul>
yum	Installed software packages	<pre>XFree86.*.log.old      /etc/pki/consumer/     cert.pem      /etc/pki/entitlement/     *.pem      /etc/pki/product/*.pem      /etc/yum/*      /etc/yum.repos.d/*      /etc/yum/     pluginconf.d/*</pre>
		<ul><li>/var/log/dnf.log</li></ul>



### Working With OSWatcher Black Box

Oracle OSWatcher Black Box (OSWbb) collects and archives operating system and network metrics that you can use to diagnose performance issues. OSWbb operates as a set of background processes on the server and gathers data on a regular basis, invoking such Unix utilities as vmstat, mpstat, netstat, iostat, and top.

OSWbb is particularly useful for Oracle RAC (Real Application Clusters) and Oracle Grid Infrastructure configurations. The RAC-DDT (Diagnostic Data Tool) script file includes OSWbb, but does not install it by default.

### **Installing OSWbb**

#### To install OSWbb:

- Log in to My Oracle Support (MOS) at https://support.oracle.com.
- Download OSWatcher from the link that is listed by Doc ID 301137.1 at https://support.oracle.com/epmos/faces/DocumentDisplay?id=301137.1.
- 3. Copy the file to the directory that you want to install OSWbb, then run the following command:

```
tar xvf oswbb{\it VERS.}tar
```

In the previous command, *VERS* represents the version number of OSWatcher, for example 832 for OSWatcher 8.32.

Extracting the tar file creates a directory named oswbb, which contains all the directories and files that are associated with OSWbb, including the startOSWbb.sh script.

4. To enable the collection of iostat information for NFS volumes, edit the OSWatcher.sh script in the oswbb directory, and set the value of nfs\_collect to 1 as follows:

```
nfs collect=1
```

### Running OSWbb

To start OSWbb, run the startOSWbb.sh script from the oswbb directory.

```
sudo ./startOSWbb.sh [frequency duration]
```

The optional frequency and duration arguments specify how often in seconds OSWbb should collect data and the number of hours for which OSWbb should run. The default values are 30 seconds and 48 hours. The following example starts OSWbb recording data at intervals of 60 seconds, and has it record data for 12 hours:

```
sudo ./startOSWbb.sh 60 12
...
Testing for discovery of OS Utilities...
VMSTAT found on your system.
```



```
IOSTAT found on your system.
MPSTAT found on your system.
IFCONFIG found on your system.
NETSTAT found on your system.
TOP found on your system.
Testing for discovery of OS CPU COUNT
oswbb is looking for the CPU COUNT on your system
CPU COUNT will be used by oswbba to automatically look for cpu problems
CPU COUNT found on your system.
CPU COUNT = 4
Discovery completed.
Starting OSWatcher Black Box v7.3.0 on date and time
With SnapshotInterval = 60
With ArchiveInterval = 12
Data is stored in directory: OSWbba archive
Starting Data Collection...
oswbb heartbeat: date and time
oswbb heartbeat: date and time + 60 seconds
```

In the previous output, *OSWbba\_archive* is the path of the archive directory that contains the OSWbb log files.

To stop OSWbb prematurely, run the stopOSWbb.sh script from the oswbb directory:

```
sudo ./stopOSWbb.sh
```

OSWbb collects data in the directories that are under the <code>oswbb/archive</code> directory, which are described in the following table.

Directory	Description
oswifconfig	Contains output from ifconfig.
oswiostat	Contains output from iostat.
oswmeminfo	Contains a listing of the contents of /proc/meminfo.
oswmpstat	Contains output from mpstat.
oswnetstat	Contains output from netstat.
oswprvtnet	If you have enable private network tracing for RAC, contains information about the status of the private networks.
oswps	Contains output from ps.
oswslabinfo	Contains a listing of the contents of /proc/slabinfo.
oswtop	Contains output from top.
oswvmstat	Contains output from vmstat.



#### OSWbb stores data in hourly archive files, which are named

system name utility name timestamp.dat. Each entry in a file is preceded by a timestamp.

### **Analyzing OSWbb Archived Files**

You can use the OSWbb analyzer (OSWbba) to provide information about system slowdowns, system hangs, and other performance problems. You can also use OSWbba to graph data that is collected from the iostat, netstat, and vmstat utilities. OSWbba requires that you have Java version 1.4.2 or a later version installed on your system.

You can download a Java RPM for Linux by visiting http://www.java.com, or you can install Java by using the dnf command:

```
sudo dnf install java-1.8.0-jdk
```

Run OSWbba from the oswbb directory as follows:

```
sudo java -jar oswbba.jar -i OSWbba archive
```

In the previous command, OSWbba\_archive is the path of the archive directory that contains the OSWbb log files.

You can use OSWbba to display the following types of performance graph:

- Process run, wait and block queues.
- CPU time spent running in system, user, and idle mode.
- Context switches and interrupts.
- Free memory and available swap.
- Reads per second, writes per second, service time for I/O requests, and percentage utilization of bandwidth for a specified block device.

You can also use OSWbba to save the analysis to a report file, which reports instances of system slowdown, spikes in run queue length, or memory shortage, describes probable causes, and offers suggestions of how to improve performance.

```
sudo java -jar oswbba.jar -i OSWbba_archive -A
```

For more information about OSWbb and OSWbba, refer to the OSWatcher Black Box User Guide (Article ID 301137.1) and the OSWatcher Black Box Analyzer User Guide (Article ID 461053.1) on My Oracle Support (MOS) at https://support.oracle.com.



### Working With Performance Co-Pilot

Performance Co-Pilot (PCP) collects operating system and network metrics that you can use to diagnose performance issues. PCP provides a monitor host that you can use to send requests for metrics and logs to a pair of collector host services that are installed on each Oracle Linux system that you monitor.

### **Installing PCP**

- Enable the ol8\_appstream and ol8\_addons yum repositories on your system.
   For more information, see Oracle<sup>®</sup> Linux: Managing Software on Oracle Linux.
- 2. Install the pcp-oracle-conf, pcp-system-tools and pcp-gui packages by using the dnf command:

```
sudo dnf install pcp-oracle-conf pcp-system-tools pcp-gui
```

3. Enable and start the Performance Metrics Collector Daemon (pmcd) and Performance Metrics Logger (pmlogger) collector host services:

```
sudo systemctl enable --now pmcd pmlogger
```

### Reviewing Information Gathered by PCP

The only metrics collected by the pmlogger service are those listed in the /var/lib/pcp/config/pmlogger/config.default file.

You can modify the frequency with which those metrics are collected in the same configuration file. For example, to increase the frequency from once every minute to once every 5 seconds:

```
# It is safe to make additions from here on ...

#
log mandatory on every 5 seconds {
  filesys.free
  filesys.used
  ...
}
```

All of the archives that the pmlogger service generates are stored in the /var/log/pcp/pmlogger/hostname directory. For more information, see the pmlogconf (1) manual page.

To verify the PCP configuration at the time that pmlogger collected specific performance metrics, use the pcp command:

```
sudo pcp -a 20220321.0.xz
```



### Using PCP Monitor Host to Analyze Performance Metrics

All of the archives that the pmlogger service generates are stored in the /var/log/pcp/pmlogger/hostname directory. If you navigate to that directory in a terminal, you can run a number of commands to review the performance metrics that have been collected.

For more information about those commands and their parameters, see their respective manual pages.

#### Review Live Performance Metrics in Real Time

To monitor all the outgoing metrics from the eth0 network interface in real time, use the pmrep command:

```
sudo pmrep -i eth0 -v network.interface.out
```

To monitor live hard drive operations for each partition with a two second interval, use the pmval command:

```
sudo pmval -t 2sec -f 3 disk.partitions.write
```

#### Review Recorded Performance Metrics

To review the data for specific performance metrics within a specified timespan, use the pmdumptext command. For example, to review resource usage metrics for CPU load, memory utilisation and disk write operations between 13:00 and 14:00 on a specific date:

```
sudo pmdumptext -Xlimu -t 10m -S @13:00 -T @14:00 \
  'kernel.all.load[1]' 'mem.util.used' 'disk.partitions.write' -a 20220321.0.xz
```

You can also use the pmstat to review system performance metrics in a format similar to that produced by the sar command. For example, to review performance metrics averaged over 10 minute interval between 09:00 and 10:00 on a specific date:

```
sudo pmstat -t 10m -S @09:00 -T @10:00 -a 20220321.0.xz
```

You can also compare the metrics between two time periods by using the pmdiff command. For example, to compare the metrics between 02:00 and 03:00 on one day to the metrics between 09:00 and 10:00 on a different day:

```
sudo pmdiff -S @02:00 -T @03:00 -B @09:00 -E @10:00 20220321.0.xz 20220320.0.xz
```

#### Review Details About Recorded Performance Metrics

To review detailed information about a specific metric, use the pminfo command. For example, to review details about free memory:

```
sudo pminfo -df mem.freemem -a 20220321.0.xz
```



### Validate System Status When Performance Metrics Were Captured

To verify the host, timezone and time period that an archive containing performance metrics contains, use the <code>pmdumplog</code> command:

sudo pmdumplog -L 20220321.0.xz

To review a list of every enabled performance metric, use the pminfo command:

sudo pminfo -a 20220321.0.xz



5

### Working With Tuned

You use the Tuned tool to monitor a system to optimize its performance under certain conditions. Tuned uses several predefined profiles to tune your system. The profiles that are provided are designed for particular use cases and fall into one of the following two categories: power-saving profiles and performance-boosting profiles. Performance-boosting profiles address low latency and high throughput for storage and the network and virtualization host performance.

You can modify the rules that are defined for each profile, as well as customize how a specific device is tuned by using a specific profile. In addition, you can configure Tuned so that any changes in device usage triggers an adjustment in the current settings so that the performance of active devices is improved and power consumption for inactive devices is reduced.

#### **About Tuned Profiles**

The following Tuned profiles are typically installed with Oracle Linux 8:

- balanced (default profile): Is a power-saving profile. This profile provides a balance between performance and power consumption. The profile uses auto-scaling and autotuning when possible. A possible drawback is increased latency.
- powersave: Is a profile that provides maximum power saving performance. The profile can minimize actual power consumption by throttling performance.

#### Note:

In some instances, the balanced profile is a better choice than the powersave profile, as it is more efficient.

- throughput-performance (default profile): Is a server profile that is optimized for high throughput. The profile disables power-savings mechanisms and enables sysctl settings to improve the throughput performance of the disk and network IO.
- latency-performance: Is a server profile that is optimized for low latency. The profile disables power-savings mechanisms and enables sysctl settings to improve latency.
- network-latency: Is a profile that provides low latency network tuning and is based on the latency-performance profile. In addition, this profile disables transparent huge pages and NUMA balancing and tunes several network-related sysctl settings.
- network-throughput: A profile for throughput network tuning. It is based on the throughput-performance profile. In addition, this profile increases kernel network buffers.
- virtual-guest (default profile): Is a profile that is designed for virtual guests and is based on the throughput-performance profile. This profile decreases virtual memory swappiness and increases disk readahead values.



- virtual-host: Is a profile that is designed for virtual hosts and is based on the throughput-performance profile. This profile decreases virtual memory swappiness, increases disk readahead values, and enables a more aggressive value of dirty pages writeback.
- desktop: Is a profile that is optimized for desktop environments and is based on the balanced profile. In addition, this profile enables scheduler autogroups for better response of interactive applications.

#### Note:

You can install additional profiles to better match your system configuration and intended use case. For example, if you are using a real-time kernel with Oracle Linux, you can use a real-time profile. These optional packages can be installed from the ol8 addons channel.

Note that real-time profiles have no effect on kernels that are not compiled with real-time support enabled.

To list all of the profiles that are currently available for installation, use the following command:

sudo dnf list tuned-profiles\*

Tuned profiles that are installed on the system by default are stored in the /usr/lib/tuned and /etc/tuned directories. Distribution-specific profiles are stored in the /usr/lib/tuned directory. Note that each profile has its own directory. Each profile directory consists of a main configuration file, tuned.conf, as well as other optional files.

If you want to use a custom profile, copy the profile directory to the /etc/tuned directory, which is the location in which custom profiles are stored. In the event there are two profiles with the same name, the custom profile that is located in /etc/tuned/ is used.

The tuned.conf file can contain one [main] section and additional sections for configuring plugin instances. Note that these sections are optional. For more information about profile configuration, see the tuned.conf (5) manual page.

#### About the Default Tuned Profiles

A default Tuned profile is automatically selected when you install Oracle Linux. The default profile that is selected is based on the given environment and the performance goals to be achieved in that particular use case. The following default profiles are provided:

- throughput-performance: Is a profile that is used in an environment that compute nodes are running Oracle Linux. This profile achieves the best throughput performance.
- virtual-guest: Is a profile that is used in an environment that virtual machines are running Oracle Linux. This profile achieves the best performance. If you are



not interested in the best performance, you can change the profile to either the balanced or powersave profile.

• balanced: Is a profile that is used for other use cases. This profile achieves balanced performance and power consumption.

### About Static and Dynamic Tuning in Tuned

Static tuning applies settings that you have defined in the configuration files for sysctl, sysfs, and other system configuration tools throughout the operating system.

You can configure the tuned service to monitor the activity of system components and dynamically tuned system settings, based on information that the service collects about the system and its current running state.

Dynamic tuning can be particularly useful in situations where you need the load on devices like the CPU, hard drives, and network adapters to consume as little power as possible when idle, but require high throughput and low latency when under a high load.

To enable dynamic tuning, set the correct value in the /etc/tuned/tuned-main.conf settings file as follows:

```
dynamic tuning = 1
```

You must then set the time interval in seconds for tuned to analyze the current system state in the same configuration file so that it can dynamically tune the system, based on the collected results, for example:

```
update_interval = 10
```

### Installing and Enabling Tuned by Using the Command Line

The following procedure describes how to install and enable Tuned, install Tuned profiles, and preset a default Tuned profile for your Oracle Linux systems.

1. If the tuned package is not already installed, install it:

```
sudo dnf install tuned
```

2. Enable and start the tuned service:

```
sudo systemctl enable --now tuned
```

3. Check the active Tuned profile:

```
sudo tuned-adm active
```

**4.** Verify that the Tuned profile is applied to the system:

```
sudo tuned-adm verify
```

The output confirms that the verification succeeded:

```
Verification succeeded, current system settings match the preset profile. See tuned log file ('/var/log/tuned/tuned.log') for details.
```

If a message about the current system settings not matching is displayed, try restarting the tuned service:

```
sudo systemctl start tuned
```



### Running Tuned in no-daemon Mode

Running tuned in no-daemon mode does not require any resident memory. However, note that when running the service in this mode, tuned does not perform any dynamic tuning. While in no-daemon mode, tuned only applies the settings and then exits.

To run tuned in no-daemon mode, you must set the following value in the /etc/tuned/tuned-main.conf settings file:

daemon = 0

#### **NOT\_SUPPORTED:**

Take note that if you decide to run tuned in no-daemon mode, be aware that some functions do not work without running the daemon. In particular, tuned no longer supports D-Bus services or the hot-plug kernel subsystem. Consequently tuned can no longer automatically roll back any settings files that were changed.

### Administering the Tuned Service and Tuned Profiles

You administer Tuned by using the tuned-adm command. The following tasks describe how to administer Tuned profiles and the tuned service on your Oracle Linux systems.

For more information, see the tuned-adm(8) and tuned(8) manual pages.

#### **Listing Tuned Profiles**

To list all of the available Tuned profiles on a system:

```
Available profiles:
- balanced - General non-specialized tuned profile
- desktop - Optimize for the desktop use-case
- latency-performance - Optimize for deterministic performance at the cost of increased power consumption
- network-latency - Optimize for deterministic performance at the cost of increased power consumption, focused on low latency network performance
- network-throughput - Optimize for streaming network throughput, generally only necessary on older CPUs or 40G+ networks
- powersave - Optimize for low power consumption
- throughput-performance - Broadly applicable tuning that provides excellent performance across a variety of common server workloads
- virtual-guest - Optimize for running inside a virtual guest
- virtual-host - Optimize for running KVM guests
Current active profile: balanced
```

Note that the current active profile is also displayed in the previous output.

To display just the currently active profile, use the following command:



sudo tuned-adm active

#### Activating a Tuned Profile

The following procedure describes how to activate a Tuned profile by using the command line. To activate a Tuned profile by using the Cockpit web console, see Oracle Linux: Use Cockpit to Set Up Performance Profiles.



To activate a Tuned profile, the tuned service must be running on your system.

Use the following command activate a specific selected Tuned profile:

sudo tuned-adm profile profile-name

To have Tuned recommend the profile that is most suitable for your system, use the tunedadm recommend command:

sudo tuned-adm recommend

The output of the command displays the recommended profile for the system on which the command is run:

virtual-guest

To activate a combination of multiple profiles, use the following command syntax:

sudo tuned-adm profile profile2

#### **Disabling Tuned**

To disable tuning temporarily, use the following command:

sudo tuned-adm off

The previous command disables any tuning settings until you restart the tuned service. When you restart the service, all of the previous tuning settings are re-applied.

You can disable tuning on a more permanent basis by stopping and disabling the tuned service as follows:

sudo systemctl disable --now tuned



6

### **Automating System Tasks**

You can use automated tasks to perform periodic backups, monitor the system, run custom scripts, as well as other administrative tasks.

The cron and anacron utilities enable you to schedule the execution of recurring tasks, referred to as *jobs*, according to a combination of the following: time, day of the month, month, day of the week, and week. With the cron utility, you can schedule jobs to run as often as every minute. If the system is down when a job is scheduled, cron does not run the job when the system restarts.

The anacron utility you to schedule a system job to run only once per day. However, if a scheduled job has not been run, that job runs when the system restarts. The anacron utility is mainly intended for use on laptop computers.

You do not usually need to run cron and anacron directly. The crond daemon executes scheduled tasks on behalf of cron and it starts anacron once every hour. crond looks in /etc/crontab or in files in /etc/cron.d for system cron job definitions, and /var/spool/cron for cron job definitions belonging to users. crond checks each job definition to see whether it should run in the current minute. If a job is scheduled for execution, crond runs it as the owner of the job definition file or, for system cron jobs, the user specified in the job definition (if any).

crond runs the <code>Oanacron</code> script in the <code>/etc/cron.hourly</code> directory as root once per hour according to the schedule in <code>/etc/cron.d/Ohourly</code>. If <code>anacron</code> is not already running and the system is connected to mains and not battery power, <code>crond</code> starts <code>anacron</code>.

anacron runs the scripts in the /etc/cron.daily, /etc/cron.weekly, and /etc/cron.monthly directories as root once per day, week or month, according to the job definitions that are scheduled in /etc/anacrontab.

### Configuring cron Jobs

System cron jobs are defined in crontab-format files in /etc/crontab or in files in /etc/cron.d. A crontab file usually consists of definitions for the SHELL, PATH, MAILTO, and HOME variables for the environment in which the jobs run, followed by the job definitions themselves. Comment lines start with a # character. Job definitions are specified in the following format:

minute hour day month day-of-week user command

Each of the fields that you can specify are defined as follows:

#### minute

Specify a value of 0-59.



#### hour

Specify a value of 0-23.

#### day

Specify a value of 1-31.

#### month

Specify a value of 1-12 or jan, feb,..., dec.

#### day-of-week

Specify a value of 0-7 (Sunday is 0 or 7) or sun, mon,...,sat.

#### user

Specify the user running the command; or, you can specify an asterisk (\*), which indicates the owner of the crontab file.

#### command

Specify the shell script or command to be run.

For the *minute* through *day-of week* fields, you can use the following special characters:

\*

Specify an asterisk (\*) for all of the valid values for the field.

-

Specify a dash (-) to indicate a range of integers, for example, 1-5.

Specify a list of values, separated by commands (,), for example, 0,2,4.

/

Specify a step value by using the forward slash (/), for example, /3 in the *hour* field. This entry is interpreted as every three hours.

For example, the following entry would run a command every five minutes on weekdays:

```
0-59/5 * * * 1-5 * command
```

Run a command at one minute past midnight on the first day of the months April, June, September, and November:

```
1 0 1 4,6,9,11 * * command
```

The root user can add job definition entries to the /etc/crontab, or add crontab-format files to the /etc/cron.d directory.



If you add an executable job script to the <code>/etc/cron.hourly</code> directory, <code>crond</code> runs the script once every hour. Your script should check that it is not already running.



For more information, see the crontab(5) manual page.

#### Controlling Access to Running cron Jobs

If permitted, users other than root can configure cron tasks by using the crontab utility. All user-defined crontab-format files are stored in the  $\sqrt{var/spool/cron}$  directory with the same name as the users that created them.

root can use the /etc/cron.allow and /etc/cron.deny files to restrict access to cron. crontab checks the access control files each time that a user tries to add or delete a cron job. If /etc/cron.allow exists, only users listed in it are allowed to use cron, and /etc/cron.deny is ignored. If /etc/cron.allow does not exist, users listed in /etc/cron.deny are not allowed to use cron. If neither file exists, only root can use cron. The format of both /etc/cron.allow and /etc/cron.deny is one user name on each line.

To create or edit a crontab file as a user, log in as that user and type the command crontab -e, which opens your crontab file in the vi editor (or the editor specified by the EDITOR or VISUAL environment variables). The file has the same format as /etc/crontab except that the user field is omitted. When you save changes to the file, these are written to the file /var/spool/cron/username. To list the contents of your crontab file, use the crontab -1 command. To delete your crontab file, use the crontab -r command.

For more information, see the crontab (5) manual page.

### Configuring anacron Jobs

System anacron jobs are defined in /etc/anacrontab, which contains definitions for the SHELL, PATH, MAILTO, RANDOM\_DELAY, and START\_HOURS\_RANGE variables for the environment in which the jobs run, followed by the job definitions themselves. Comment lines start with a # character.

RANDOM\_DELAY is the maximum number of random time in minutes that anacron adds to the delay parameter for a job. The default minimum delay is 6 minutes. The random offset is intended to prevent anacron overloading the system with too many jobs at the same time.

START\_HOURS\_RANGE is the time range of hours during the day when anacron can run scheduled jobs.

Job definitions are specified in the following format:

```
period delay job-id command
```

The following is a description of the fields that may be included:

#### period

Frequency of job execution specified in days or as @daily, @weekly, or @monthly for once per day, week, or month.

#### delay

Number of minutes to wait before running a job.

#### job-id

Unique name for the job in log files.



#### command

The shell script or command to be run.

The following entries are taken from the default /etc/anacrontab file:

```
SHELL=/bin/sh
PATH=/sbin:/bin:/usr/sbin:/usr/bin
MAILTO=root
# the maximal random delay added to the base delay of the jobs
RANDOM DELAY=45
# the jobs will be started during the following hours only
START HOURS RANGE=3-22
#period in days delay in minutes job-identifier command
             5
                             cron.daily nice run-parts /etc/
cron.daily
             25
7
                             cron.weekly nice run-parts /etc/
cron.weekly
@monthly
           45
                               cron.monthly nice run-parts /etc/
cron.monthly
```

By default, anacron runs jobs between 03:00 and 22:00 and randomly delays jobs by between 11 and 50 minutes. The job scripts in /etc/cron.daily, run anywhere between 03:11 and 03:50 every day if the system is running, or after the system is booted and the time is less than 22:00. The run-parts script sequentially executes every program within the directory specified as its argument.

Scripts in /etc/cron.weekly run once per week with a delay offset of between 31 and 70 minutes.

Scripts in /etc/cron.monthly run once per week with a delay offset of between 51 and 90 minutes.

For more information, see the anacron(8) and anacrontab(5) manual pages.

# **Running One-Time Tasks**

You can use the at command to schedule a one-time task to run at a specified time, or the batch command to schedule a one-time task to run when the system load average drops below 0.8. The atd service must be running to use at or batch.

```
sudo systemctl is-active atd
```

at takes a time as its argument and reads the commands to be run from the standard input. For example, run the commands in the file atjob in 20 minutes time:

```
at now + 20 minutes < ./atjob
job 1 at 2013-03-19 11:25
```

The atq command shows the at jobs that are queued to run, for example:

```
sudo atq
1 2013-03-19 11:25 a root
```

The batch command also reads command from the standard input, but it does not run until the system load average drops below 0.8, for example:



```
sudo batch < batchjob
job 2 at 2013-03-19 11:31</pre>
```

To cancel one or more queued jobs, specify their job numbers to the atrm command, for example:

```
sudo atrm 1 2
```

For more information, see the at (1) manual page.

### Changing the Behavior of Batch Jobs

The load average of a system, as displayed by the uptime and w commands, represents the average number of processes that are queued to run on the CPUs or CPU cores over a given time period. Typically, a system might not considered overloaded until the load average exceeds 0.8 times the number of CPUs or CPU cores. On such systems, you would usually want atd to be able to run batch jobs when the load average drops below the number of CPUs or CPU cores, rather than the default limit of 0.8. For example, on a system with 4 CPU cores, you could set the load-average limit above which atd will not run batch jobs to 3.2.

If you know that a batch job typically takes more than a minute to run, you can also change the minimum interval that atd waits between starting batch jobs. The default minimum interval is 60 seconds.

To change the load-average limit and minimum interval time for batch jobs:

1. Edit the atd configuration file, /etc/sysconfig/atd, uncomment the line that defines the OPTS variable, and edit the line to specify the new load-average limit and minimum interval time. for example:

```
OPTS="-b 100 -1 3"
```

This example sets the minimum interval to 100 seconds and the load-average limit to 3.

2. Restart the atd service:

```
sudo systemctl restart atd
```

3. Verify that the atd daemon is running with the new minimum interval and load-average limit, for example:

For more information, see the <code>systemctl(1)</code> and <code>atd(8)</code> manual pages.



7

# **Configuring and Using Auditing**

Auditing collects data at the kernel level that you can then analyze to identify unauthorized activity. Auditing collects data in greater detail than system logging does; but, note that most audited events are uninteresting and insignificant. Because the process of examining audit trails to locate events of interest can be a significantly challenging, you may consider automating this process.

The audit configuration file, /etc/audit/auditd.conf, defines the following:

- Data retention policy
- Maximum size of the audit volume
- Action to take if the capacity of the audit volume is exceeded
- Locations of local and remote audit trail volumes

The default audit trail volume is /var/log/audit/audit.log. See the auditd.conf(5) manual page for more information.

By default, auditing captures specific events such as system logins, modifications to accounts, and <code>sudo</code> actions. You can alsoA configure auditing to capture detailed system call activity or modifications to certain files. The kernel audit daemon (<code>auditd</code>) records the events that you configure, including the event type, a time stamp, the associated user ID, and success or failure of the system call.

The entries in the audit rules file, /etc/audit/audit.rules, determine which events are audited. Each rule is a command-line option that is passed to the auditctl command. You should typically configure this file to match your site's security policy.

The following are examples of rules that you might set in the /etc/audit/audit.rules file:

Record all unsuccessful exits from open and truncate system calls for files in the /etc directory hierarchy.

```
-a exit, always -S open -S truncate -F /etc -F success=0
```

To record all files opened by a user with UID 10:

```
-a exit, always -S open -F uid=10
```

To record all files that have been written to or files with their attributes changed by any user who originally logged in with a UID of 500 or greater:

```
-a exit, always -S open -F auid>=500 -F perm=wa
```

To record requests for write or file attribute change access to the /etc/sudoers file and tag such a record with the string sudoers-change:

```
-w /etc/sudoers -p wa -k sudoers-change
```

To record requests for write and file attribute change access to the /etc directory hierarchy:

```
-w /etc/ -p wa
```



To require a reboot after changing the audit configuration:

-e 2

If specified, this rule should appear at the end of the /etc/audit/audit.rules file.

For more examples of audit rules, refer to the /usr/share/doc/audit-version/stig.rules file. See also the auditctl(8) and audit.rules(7) manual pages.

Stringent auditing requirements can impose a significant performance overhead and generate large amounts of audit data. Some site security policies stipulate that a system must shut down if events cannot be recorded because the audit volumes have exceeded their capacity. As a general rule, you should direct audit data to separate file systems in rotation to prevent overspill and to facilitate backups.

You can use the -k option to tag audit records so that you can locate them more easily in an audit volume with the ausearch command. For example, to examine records tagged with the string sudoers-change, you would enter:

```
sudo ausearch -k sudoers-change
```

The aureport command generates summaries of audit data. You can set up cron jobs that run aureport periodically to generate reports of interest. For example, the following command generates a reports that shows every login event from 1 second after midnight on the previous day until the current time:

```
sudo aureport -l -i -ts yesterday -te now
```

See the ausearch (8) and aureport (8) manual pages for more information.

See https://docs.oracle.com/en/learn/ol-auditd/ for a hands-on tutorial on using the auditing tools on Oracle Linux.

# Working With System Log files

The log files contain messages about the system, kernel, services, and applications. The <code>journald</code> logging daemon, which is part of <code>systemd</code>, records system messages in non-persistent journal files in memory and in the <code>/run/log/journal</code> directory. <code>journald</code> forwards messages to the system logging daemon, <code>rsyslog</code>. As files in <code>/run</code> are volatile, the log data is lost after a reboot unless you create the directory <code>/var/log/journal</code>. You can use the <code>journalctl</code> command to query the journal logs.

For more information, see the journalctl(1) and systemd-journald.service(8) manual pages.

See https://docs.oracle.com/en/learn/system\_logging\_linux8/ for a hands-on tutorial introducing system logging tools.

## About Logging Configuration (/etc/rsyslog.conf)

The configuration file for rsyslogd is /etc/rsyslog.conf, which contains global directives, module directives, and rules. By default, rsyslog processes and archives only syslog messages. If required, you can configure rsyslog to archive any other messages that journald forwards, including kernel, boot, initrd, stdout, and stderr messages.



Global directives specify configuration options that apply to the rsyslogd daemon. All configuration directives must start with a dollar sign (\$) and only one directive can be specified on each line. The following example specifies the maximum size of the rsyslog message queue:

\$MainMsqQueueSize 50000

The available configuration directives are described in the file /usr/share/doc/rsyslog-version-number/rsyslog conf global.html.

The design of rsyslog ensures that functionality is dynamically loaded from modules, which provide configuration directives. To load a module, specify the following directive:

\$ModLoad MODULE name

Modules have the following main categories:

- Input modules gather messages from various sources. Input module names always start with the im prefix (examples include imfile and imrelp).
- Filter modules enable rsyslogd to filter messages according to specified rules. The name of a filter module always starts with the fm prefix.
- Library modules provide functionality for other loadable modules. rsyslogd loads library modules automatically when required. You cannot configure the loading of library modules.
- Output modules provide the facility to store messages in a database or on other servers in a network, or to encrypt them. Output module names always starts with the om prefix (examples include omsnmp and omrelp).
- Message modification modules change the content of an rsyslog message.
- Parser modules enable rsyslogd to parse the message content of messages that it receives. The name of a parser module always starts with the pm prefix.
- String generator modules generate strings based on the content of messages in cooperation with rsyslog's template feature. The name of a string generator module always starts with the sm prefix.

Input modules receive messages, which pass them to one or more parser modules. A parser module creates a representation of a message in memory, possibly modifying the message, and passes the internal representation to output modules, which can also modify the content before outputting the message.

A description of the available modules can be found at https://www.rsyslog.com/doc/rsyslog\_conf\_modules.html.

An rsyslog rule consists of a filter part, which selects a subset of messages, and an action part, which specifies what to do with the selected messages. To define a rule in the /etc/rsyslog.conf configuration file, specify a filter and an action on a single line, separated by one or more tabs or spaces.

You can configure rsyslog to filter messages according to various properties. The following are the most commonly used filters:

• Expression-based filters, written in the rsyslog scripting language, select messages according to arithmetic, boolean, or string values.



- Facility/priority-based filters filter messages based on facility and priority values that take the form <code>facility.priority</code>.
- Property-based filters filter messages by properties such as timegenerated or syslogtag.

The following table describes the available facility keywords for facility or priority-based filters.

Facility Keyword	Description	
auth, authpriv	Security, authentication, or authorization messages.	
cron	crond messages.	
daemon	Messages from system daemons other than crond and rsyslogd.	
kern	Kernel messages.	
lpr	Line printer subsystem.	
mail	Mail system.	
news	Network news subsystem.	
syslog	Messages generated internally by rsyslogd.	
user	User-level messages.	
UUCP	UUCP subsystem.	
local0-local7	Local use.	

The following table describes the available priority keywords for facility or priority-based filters, in ascending order of importance.

Priority Keyword	Description	
debug	Debug-level messages.	
info	Informational messages.	
notice	Normal but significant condition.	
warning	Warning conditions.	
err	Error conditions.	
crit	Critical conditions.	
alert	Immediate action required.	
emerg	System is unstable.	

All messages of the specified priority and higher are logged according to the specified action. An asterisk (\*) wildcard specifies all facilities or priorities. Separate the names of multiple facilities and priorities on a line with commas (,). Separate multiple filters on one line with semicolons (;). Precede a priority with an exclamation mark (!) to select all messages except those with that priority.

The following are examples of facility/priority-based filters.



Select all kernel messages with any priority as follows:

kern.\*

Select all mail messages with crit or higher priority as follows:

mail.crit

Select all daemon and kern messages with warning or err priority as follows:

daemon, kern.warning, err

Select all cron messages except those with info or debug priority as follows:

cron.!info,!debug

By default, /etc/rsyslog.conf includes the following rules:

```
# Log all kernel messages to the console.
# Logging much else clutters up the screen.
#kern.*
                                                         /dev/console
# Log anything (except mail) of level info or higher.
# Don't log private authentication messages!
*.info;mail.none;authpriv.none;cron.none
                                                         /var/log/messages
# The authpriv file has restricted access.
authpriv.*
                                                         /var/log/secure
# Log all the mail messages in one place.
mail.*
                                                         -/var/log/maillog
# Log cron stuff
cron.*
                                                         /var/log/cron
# Everybody gets emergency messages
# Save news errors of level crit and higher in a special file.
uucp, news.crit
                                                         /var/log/spooler
# Save boot messages also to boot.log
local7.*
                                                         /var/log/boot.log
```

You can send the logs to a central log server over TCP by adding the following entry to the forwarding rules section of /etc/rsyslog.conf on each log client:

```
*.* @@logsvr:port
```

In the previous example, *logsvr* is the domain name or IP address of the log server and port is the port number (usually, 514).

On the log server, add the following entry to the MODULES section of /etc/rsyslog.conf:

```
$ModLoad imtcp
$InputTCPServerRun port
```

In the previous example, *port* corresponds to the port number that you set on the log clients.

To manage the rotation and archival of the correct logs, edit /etc/logrotate.d/ syslog so that it references each of the log files that are defined in the RULES section of /etc/rsyslog.conf. You can configure how often the logs are rotated and how many past copies of the logs are archived by editing /etc/logrotate.conf.

It is recommended that you configure Logwatch on your log server to monitor the logs for suspicious messages, and disable Logwatch on log clients. However, if you do use Logwatch, disable high precision timestamps by adding the following entry to the GLOBAL DIRECTIVES section of /etc/rsyslog.conf on each system:

\$ActionFileDefaultTemplate RSYSLOG TraditionalFileFormat

For more information, see the <code>logrotate(8)</code>, <code>logwatch(8)</code>, <code>rsyslogd(8)</code> and <code>rsyslog.conf(5)</code> manual pages. See also the HTML documentation in the <code>/usr/share/doc/rsyslog-5.8.10</code> directory and the documentation at <a href="https://www.rsyslog.com/doc/manual.html">https://www.rsyslog.com/doc/manual.html</a>.

### **Configuring Logwatch**

Logwatch is a monitoring system that you can configure to report on areas of interest in the system logs. After you install the <code>logwatch</code> package, the <code>/etc/cron.daily/Ologwatch</code> script runs every night and sends an email report to <code>root</code>. You can set local configuration options in <code>/etc/logwatch/conf/logwatch.conf</code> that override the main configuration file <code>/usr/share/logwatch/default.conf/logwatch.conf</code>, including the following:

- Log files to monitor, including log files that are stored for other hosts.
- Names of services to monitor, or to be excluded from monitoring.
- Level of detail to report.
- User to be sent an emailed report.

You can also run logwatch directly from the command line.

For more information, see the logwatch (8) manual page.

## **Using Process Accounting**

The psacct package implements the process accounting service in addition to the following utilities that you can use to monitor process activities:

#### ac

Displays connection times in hours for a user as recorded in the wtmp file (by default, /var/log/wtmp).

### accton

Turns on process accounting to the specified file. If you do not specify a file name argument, process accounting is stopped. The default system accounting file is /var/account/pacct.

#### lastcomm

Displays information about previously executed commands as recorded in the system accounting file.



### sa

Summarizes information about previously executed commands as recorded in the system accounting file.



As for any logging activity, ensure that the file system has enough space to store the system accounting and wtmp files. Monitor the size of the files and, if necessary, truncate them.

For more information, see the ac(1), accton(8), lastcomm(1), and sa(8) manual pages.



# Working With Kernel Dumps

The Kdump feature provides a kernel crash dumping mechanism in Oracle Linux. The kdump service enables you to save the contents of the system's memory for later analysis. The second kernel resides in a reserved part of the system memory.

Kdump uses the kexec system call to boot into the second kernel, called a *capture kernel*, without the need to reboot the system, and then captures the contents of the crashed kernel's memory as a crash dump (vmcore) and saves it. The vmcore crash dump can help with determining the cause of the crash.

Oracle recommends that you enable the Kdump feature because a crash dump might be the only information that is available if a system failure occurs. Ensuring that Kdump is enabled is critical in many mission-critical environments.

Prior to enabling Kdump, ensure that your system meets all of the memory requirements for using Kdump. To capture a kernel crash dump and save it for further analysis, you must permanently reserve part of the system's memory for the capture kernel. Note that this part of the system's memory will no longer be available to the main kernel. The following table lists the minimum amount of reserved memory that is required to use Kdump, based on the system's architecture and the amount of available memory.

Architecture	Available Memory	Minimum Reserved Memory
x86_64	1 GB to 64 GB	160 MB of RAM
	64 GB to 1 TB	256 MB of RAM
	1 TB and more	512 MB of RAM
Arm (aarch64)	2 TB and more	512 MB of RAM

# Kdump Installation and Configuration

This section describes how to install and configure Kdump by using the command line.

For information about configuring Kdump by using the Cockpit web console, see Oracle Linux: Use Cockpit to Configure Kdump

### Files That Are Used by Kdump

When you install and configure Kdump, the following files are modified:

/boot/grub2/grub.cfg

Appends the crashkernel option to the kernel line to specify the amount of reserved memory and any offset value.



### /etc/kdump.conf

Sets the location in which the dump file can be written, the filtering level for the makedumpfile command, and the default behavior to take if the dump fails. See the comments in the file for information about the supported parameters.

When you edit these files, you must reboot the system for the changes to take effect.

For more information, see the kdump.conf(5) manual page.

## Installing and Configuring Kdump

During an Oracle Linux interactive installation with the graphical installer, you have the option to enable Kdump and specify how much system memory is reserved for Kdump. The installer screen is titled **Kdump** and is available from the main Installation Summary screen of the installer.

If you do not enable Kdump at installation time, or it is not enabled by default during an installation, as in the case of a custom kickstart installation, you can install and enable the feature by using the command line.

Before you install and configure Kdump by using the command line, ensure that your system meets all of the necessary memory requirements. For details, see Working With Kernel Dumps.

1. If the kdump package is not already installed on your system, install it:

```
sudo dnf install kexec-tools
```

2. As the root user, edit the /etc/default/grub file and set the crashkernel= option to the required value.

For example, you would reserve 64 MB of memory as follows:

```
crashkernel=64M
```

You can also set the amount of reserved memory as a variable by using the following syntax: crashkernel=range1:size1, range2:size2

For example, you might set the memory as a variable as follows:

```
crashkernel=512M-2G:64M,2G-:128M
```

3. (Optional) If necessary, offset the reserved memory.

Because the crashkernel reservation occurs very early, some systems require that you reserve memory with a certain fixed offset. When a fixed offset is specified, the reserved memory begins at that point. For example, you would reserve 128 MB of memory, starting at 16 MB as follows:

```
crashkernel=128M@16M
```

Note that if no offset parameter is set, Kdump offsets reserved memory automatically.

4. Refresh the grub configuration to apply your changes:

```
sudo grub2-mkconfig -o /boot/grub2/grub.cfg
```

5. Reboot the system and finish configuring Kdump.

For instructions, see Configuring the Kdump Output Location.

6. When you have finished configuring Kdump, enable the kdump service:

sudo systemctl enable --now kdump.service

## Configuring the Kdump Output Location

After installing Kdump, you can define the location in which the resulting output should be saved. For Oracle Linux, Kdump files are stored in the <code>/var/crash</code> directory by default.

To save the result to other locations, such as NFS mounts, externally mounted drives, and remote file servers, edit the /etc/kdump.conf file and remove the # comment character at the beginning of each line that you want to enable.

For example, to add a new directory location, prefix it with the path keyword:

```
path /usr/local/cores
```

Use raw to output directly to a specific device in the /dev directory. You can also manually specify the output file system for a particular device by using its label, name or UUID, for example:

```
ext4 UUID=5b065be6-9ce0-4154-8bf3-b7c4c7dc7365
```

Kdump files can also be transferred over a secure shell connection, as shown in the following example:

```
ssh user@example.com
sshkey /root/.ssh/mykey
```

It is also possible to export the result to a compatible network share:

```
nfs example.com:/output
```

When you have finished configuring the output location for Kdump, enable the kdump service as follows:

```
sudo systemctl enable --now kdump.service
```

## Configuring the Default Kdump Failure State

By default, if kdump fails to output its result to the configured outlook locations, it reboots the server. This action deletes any data that has been collected for the dump, so you should uncomment and change the default value in the /etc/kdump.conf file as follows:

```
default dump to rootfs
```

The <code>dump\_to\_rootfs</code> option attempts to save the result to a local directory, which can be particularly useful if a network share is unreachable. Using <code>shell</code> instead enables you to copy the data manually from the command line.



The poweroff, restart, and haltoptions are also valid for the default kdump failure state; but, you will lose the collected data if those actions are performed.



# **Analyzing Kdump Output**

The crash utility provides a shell prompt that enables you to analyze the contents of your kdump core dumps, which is particularly useful when troubleshooting problems.

1. If the crash package is not installed on your system, install it:

```
sudo dnf install crash
```

2. Identify the currently running kernel, for example:

```
uname -r
```

The output of the previous command is similar to the following:

```
4.18.0-80.el8.x86 64
```

3. Provide the location of the kernel debuginfo module and the location of the core dump as parameters to the crash utility, for example:

```
sudo crash /usr/lib/debug/lib/modules/4.18.0-80.el8.x86_64/vmlinux \
   /var/crash/127.0.0.1-2019-10-28-12:38:25/vmcore
```

In the previous command, 4.18.0-80.el8.x86\_64 is the currently running kernel and 127.0.0.1-2019-10-28-12:38:25 represents the *ipaddress-timestamp*.

4. Inside the crash shell, you can use the help log command to better understand how to use the log command.

You can also use the bt, ps, vm, and files commands to get more information about the core dump.

5. When you have finished analyzing the core dump, exit the shell.

For more detailed information about using the crash utility, see the crash (8) manual page.

# **Using Early Kdump**

New in Oracle Linux 8, early Kdump enables the crash kernel and initramfs to load early enough to capture vmcore information for early crashes.

Because the kdump service starts too late, early crashes are not able to use normal kdump kernel booting. As a result, information about early crashes is lost. To address this issue, you can enable early Kdump by adding a dracut module so that the crash kernel and initramfs are loaded as early as possible. When early Kdump is enabled, the files are loaded just like a normal kdump, which is disabled by default.

Note that early Kdump does not support Fadump currently.

For more information about configuring early Kdump on your Oracle Linux 8 systems, see the step-by-step instructions in the /usr/share/doc/kexec-tools/early-kdump-howto.txt file.



# Using Kdump with OCFS2

By default, a fenced node in an OCFS2 cluster restarts instead of panicking so that it can quickly rejoin the cluster. If the reason for the restart is not apparent, you can change the node's behavior so that it panics and generates a vmcore for analysis.

To configure a node to panic when it next fences, run the following command on the node after the cluster starts:

echo panic > /sys/kernel/config/cluster/cluster\_name/fence\_method

In the previous command, <code>cluster\_name</code> is the name of the cluster. To set the value after each reboot of the system, add this line to /etc/rc.local. To restore the default behavior, set the value of <code>fence\_method</code> to <code>reset</code> instead of! <code>panic</code> and remove the line from /etc/rc.local.

For more information, see Oracle® Linux 8: Setting Up High Availability Clustering.

